NAME

ovn-nbctl - Open Virtual Network northbound db management utility

SYNOPSIS

ovn-nbctl [options] command [arg...]

DESCRIPTION

This utility can be used to manage the OVN northbound database.

GENERAL COMMANDS

init Initializes the database, if it is empty. If the database has already been initialized, this command has no effect.

show [switch | router]

Prints a brief overview of the database contents. If *switch* is provided, only records related to that logical switch are shown. If *router* is provided, only records related to that logical router are shown.

LOGICAL SWITCH COMMANDS

ls-add Creates a new, unnamed logical switch, which initially has no ports. The switch does not have a name, other commands must refer to this switch by its UUID.

[--may-exist | --add-duplicate] ls-add switch

Creates a new logical switch named switch, which initially has no ports.

The OVN northbound database schema does not require logical switch names to be unique, but the whole point to the names is to provide an easy way for humans to refer to the switches, making duplicate names unhelpful. Thus, without any options, this command regards it as an error if *switch* is a duplicate name. With —**may**—**exist**, adding a duplicate name succeeds but does not create a new logical switch. With —**add**—**duplicate**, the command really creates a new logical switch with a duplicate name. It is an error to specify both options. If there are multiple logical switches with a duplicate name, configure the logical switches using the UUID instead of the *switch* name.

[--if-exists] ls-del switch

Deletes switch. It is an error if switch does not exist, unless --if-exists is specified.

ls-list Lists all existing switches on standard output, one per line.

LOGICAL SWITCH ACL COMMANDS

[--log] [--may-exist] acl-add switch direction priority match action

Adds the specified ACL to *switch. direction* must be either **from-lport** or **to-lport**. *priority* must be between **0** and **32767**, inclusive. If **--log** is specified, packet logging is enabled for the ACL. A full description of the fields are in **ovn-nb**(5). If **--may-exist** is specified, adding a duplicated ACL succeeds but the ACL is not really created. Without **--may-exist**, adding a duplicated ACL results in error.

acl-del switch [direction [priority match]]

Deletes ACLs from *switch*. If only *switch* is supplied, all the ACLs from the logical switch are deleted. If *direction* is also specified, then all the flows in that direction will be deleted from the logical switch. If all the fields are given, then a single flow that matches all the fields will be deleted.

acl-list switch

Lists the ACLs on switch.

LOGICAL SWITCH PORT COMMANDS

[--may-exist] lsp-add switch port

Creates on *lswitch* a new logical switch port named *port*.

It is an error if a logical port named *port* already exists, unless —**may**—**exist** is specified. Regardless of —**may**—**exist**, it is an error if the existing port is in some logical switch other than *switch* or if it has a parent port.

[--may-exist] lsp-add switch port parent tag_request

Creates on *switch* a logical switch port named *port* that is a child of *parent* that is identified with VLAN ID *tag_request*, which must be between **0** and **4095**, inclusive. If *tag_request* is **0**, **ovn—northd** generates a tag that is unique in the scope of *parent*. This is useful in cases such as virtualized container environments where Open vSwitch does not have a direct connection to the container's port and it must be shared with the virtual machine's port.

It is an error if a logical port named *port* already exists, unless **—may—exist** is specified. Regardless of **—may—exist**, it is an error if the existing port is not in *switch* or if it does not have the specified *parent* and *tag_request*.

[--if-exists] lsp-del port

Deletes port. It is an error if port does not exist, unless --if-exists is specified.

lsp-list switch

Lists all the logical switch ports within *switch* on standard output, one per line.

lsp-get-parent port

If set, get the parent port of *port*. If not set, print nothing.

lsp-get-tag port

If set, get the tag for *port* traffic. If not set, print nothing.

lsp-set-addresses *port* [address]...

Sets the addresses associated with port to address. Each address should be one of the following:

an Ethernet address, optionally followed by a space and one or more IP addresses

OVN delivers packets for the Ethernet address to this port.

unknown

OVN delivers unicast Ethernet packets whose destination MAC address is not in any logical port's addresses column to ports with address **unknown**.

dvnamic

Use this keyword to make **ovn-northd** generate a globally unique MAC address and choose an unused IPv4 address with the logical port's subnet and store them in the port's **dynamic_addresses** column.

router Accepted only when the **type** of the logical switch port is **router**. This indicates that the Ethernet, IPv4, and IPv6 addresses for this logical switch port should be obtained from the connected logical router port, as specified by **router-port** in **lsp-set-options**.

Multiple addresses may be set. If no *address* argument is given, *port* will have no addresses associated with it.

lsp-get-addresses port

Lists all the addresses associated with *port* on standard output, one per line.

lsp-set-port-security *port* [addrs]...

Sets the port security addresses associated with *port* to *addrs*. Multiple sets of addresses may be set by using multiple *addrs* arguments. If no *addrs* argument is given, *port* will not have port security enabled.

Port security limits the addresses from which a logical port may send packets and to which it may receive packets. See the **ovn-nb**(5) documentation for the **port_security** column in the **Logical_Switch_Port** table for details.

lsp-get-port-security port

Lists all the port security addresses associated with *port* on standard output, one per line.

lsp-get-up port

Prints the state of *port*, either **up** or **down**.

lsp-set-enabled port state

Set the administrative state of *port*, either **enabled** or **disabled**. When a port is disabled, no traffic is allowed into or out of the port.

lsp-get-enabled port

Prints the administrative state of *port*, either **enabled** or **disabled**.

lsp-set-type *port type*

Set the type for the logical port. No special types have been implemented yet.

lsp-get-type port

Get the type for the logical port.

lsp-set-options *port* [*key=value*]...

Set type-specific key-value options for the logical port.

lsp-get-options port

Get the type-specific options for the logical port.

LOGICAL ROUTER COMMANDS

lr-add Creates a new, unnamed logical router, which initially has no ports. The router does not have a name, other commands must refer to this router by its UUID.

[--may-exist | --add-duplicate] lr-add router

Creates a new logical router named *router*, which initially has no ports.

The OVN northbound database schema does not require logical router names to be unique, but the whole point to the names is to provide an easy way for humans to refer to the routers, making duplicate names unhelpful. Thus, without any options, this command regards it as an error if *router* is a duplicate name. With ——may—exist, adding a duplicate name succeeds but does not create a new logical router. With ——add—duplicate, the command really creates a new logical router with a duplicate name. It is an error to specify both options. If there are multiple logical routers with a duplicate name, configure the logical routers using the UUID instead of the *router* name.

[--if-exists] lr-del router

Deletes *router*. It is an error if *router* does not exist, unless —**if**—**exists** is specified.

lr-list Lists all existing routers on standard output, one per line.

LOGICAL ROUTER PORT COMMANDS

[--may-exist] lrp-add router port mac network... [peer=peer]

Creates on *router* a new logical router port named *port* with Ethernet address *mac* and one or more IP address/netmask for each *network*.

The optional argument **peer** identifies a logical router port that connects to this one. The following example adds a router port with an IPv4 and IPv6 address with peer **lr1**:

lrp-add lr0 lrp0 00:11:22:33:44:55 192.168.0.1/24 2001:db8::1/64 peer=lr1

It is an error if a logical router port named *port* already exists, unless ——may—exist is specified. Regardless of —may—exist, it is an error if the existing router port is in some logical router other than *router*.

[--if-exists] lrp-del port

Deletes *port*. It is an error if *port* does not exist, unless —**if**—**exists** is specified.

lrp-list router

Lists all the logical router ports within *router* on standard output, one per line.

lrp-set-enabled *port state*

Set the administrative state of *port*, either **enabled** or **disabled**. When a port is disabled, no traffic is allowed into or out of the port.

lrp-get-enabled port

Prints the administrative state of *port*, either **enabled** or **disabled**.

LOGICAL ROUTER STATIC ROUTE COMMANDS

[--may-exist] [--policy=POLICY] lr-route-add router prefix nexthop [port]

Adds the specified route to *router*. *prefix* describes an IPv4 or IPv6 prefix for this route, such as **192.168.100.0/24**. *nexthop* specifies the gateway to use for this route, which should be the IP address of one of *router* logical router ports or the IP address of a logical port. If *port* is specified, packets that match this route will be sent out that port. When *port* is omitted, OVN infers the output port based on *nexthop*.

—policy describes the policy used to make routing decisions. This should be one of "dst-ip" or "src-ip". If not specified, the default is "dst-ip".

It is an error if a route with *prefix* already exists, unless — **may**—**exist** is specified.

[--if-exists] lr-route-del router [prefix]

Deletes routes from *router*. If only *router* is supplied, all the routes from the logical router are deleted. If *prefix* is also specified, then all the routes that match the prefix will be deleted from the logical router.

It is an error if *prefix* is specified and there is no matching route entry, unless —**if-exists** is specified.

lr-route-list *router*

Lists the routes on *router*.

NAT COMMANDS

[--may-exist] lr-nat-add router type external_ip logical_ip [logical_port external_mac]

Adds the specified NAT to *router*. The *type* must be one of **snat**, **dnat**, or **dnat_and_snat**. The *external_ip* is an IPv4 address. The *logical_ip* is an IPv4 network (e.g 192.168.1.0/24) or an IPv4 address. The *logical_port* and *external_mac* are only accepted when *router* is a distributed router (rather than a gateway router) and *type* is **dnat_and_snat**. The *logical_port* is the name of an existing logical switch port where the *logical_ip* resides. The *external_mac* is an Ethernet address.

When *type* is **dnat**, the externally visible IP address *external_ip* is DNATted to the IP address *logical_ip* in the logical space.

When type is **snat**, IP packets with their source IP address that either matches the IP address in logical_ip or is in the network provided by logical_ip is SNATed into the IP address in external_ip.

When *type* is **dnat_and_snat**, the externally visible IP address *external_ip* is DNATted to the IP address *logical_ip* in the logical space. In addition, IP packets with the source IP address that matches *logical_ip* is SNATed into the IP address in *external_ip*.

When the *logical_port* and *external_mac* are specified, the NAT rule will be programmed on the chassis where the *logical_port* resides. This includes ARP replies for the *external_ip*, which return the value of *external_mac*. All packets transmitted with source IP address equal to *external_ip* will be sent using the *external_mac*.

It is an error if a NAT already exists with the same values of *router*, *type*, *external_ip*, and *logical_ip*, unless —**may**—**exist** is specified. When —**may**—**exist**, *logical_port*, and *external_mac* are all specified, the existing values of *logical_port* and *external_mac* are overwritten.

[--if-exists] lr-nat-del router [type [ip]]

Deletes NATs from *router*. If only *router* is supplied, all the NATs from the logical router are deleted. If *type* is also specified, then all the NATs that match the *type* will be deleted from the logical router. If all the fields are given, then a single NAT rule that matches all the fields will be deleted. When *type* is **snat**, the *ip* should be logical_ip. When *type* is **dnat** or **dnat_and_snat**, the *ip* should be external_ip.

It is an error if *ip* is specified and there is no matching NAT entry, unless —**if**—**exists** is specified.

lr-nat-list router

Lists the NATs on router.

LOAD BALANCER COMMANDS

[--may-exist | --add-duplicate] lb-add lb vip ips [protocol]

Creates a new load balancer named *lb* with the provided *vip* and *ips* or adds the *vip* to an existing *lb*. *vip* should be a virtual IPv4 address (or an IPv4 address and a port number with: as a separator). Examples for *vip* are **192.168.1.4** and **192.168.1.5:8080**. *ips* should be comma separated IPv4 endpoints (or comma separated IPv4 addresses and port numbers with: as a separator). Examples for *ips* are **10.0.0.1,10.0.0.2** or **20.0.0.10:8800,20.0.0.11:8800**.

The optional argument *protocol* must be either **tcp** or **udp**. This argument is useful when a port number is provided as part of the *vip*. If the *protocol* is unspecified and a port number is provided as part of the *vip*, OVN assumes the *protocol* to be **tcp**.

It is an error if the *vip* already exists in the load balancer named *lb*, unless —**may-exist** is specified. With —**add-duplicate**, the command really creates a new load balancer with a duplicate name.

The following example adds a load balancer.

lb-add lb0 30.0.0.10:80 192.168.10.10:80,192.168.10.20:80,192.168.10.30:80 udp

[--if-exists] lb-del lb [vip]

Deletes *lb* or the *vip* from *lb*. If *vip* is supplied, only the *vip* will be deleted from the *lb*. If only the *lb* is supplied, the *lb* will be deleted. It is an error if *vip* does not already exist in *lb*, unless —if—exists is specified.

lb-**list** [*lb*]

Lists the LBs. If *lb* is also specified, then only the specified *lb* will be listed.

[--may-exist] ls-lb-add switch lb

Adds the specified *lb* to *switch*. It is an error if a load balancer named *lb* already exists in the *switch*, unless —**may**—**exist** is specified.

[--if-exists] ls-lb-del switch [lb]

Removes *lb* from *switch*. If only *switch* is supplied, all the LBs from the logical switch are removed. If *lb* is also specified, then only the *lb* will be removed from the logical switch. It is an error if *lb* does not exist in the *switch*, unless —**if**—**exists** is specified.

ls-lb-list switch

Lists the LBs for the given switch.

[--may-exist] lr-lb-add router lb

Adds the specified *lb* to *router*. It is an error if a load balancer named *lb* already exists in the *router*, unless —**may**—**exist** is specified.

[--if-exists] lr-lb-del router [lb]

Removes *lb* from *router*. If only *router* is supplied, all the LBs from the logical router are removed. If *lb* is also specified, then only the *lb* will be removed from the logical router. It is an error if *lb* does not exist in the *router*, unless —**if**—**exists** is specified.

lr-lb-list router

Lists the LBs for the given *router*.

DHCP OPTIONS COMMANDS

dhcp-options-create *cidr* [*key=value*]

Creates a new DHCP Options entry in the **DHCP_Options** table with the specified **cidr** and optional **external-ids**.

dhcp-options-list

Lists the DHCP Options entries.

${\bf dhcp\text{-}options\text{-}del}\ \textit{dhcp-}option$

Deletes the DHCP Options entry referred by dhcp-option UUID.

dhcp-options-set-options *dhcp-option* [*key=value*]...

Set the DHCP Options for the *dhcp-option* UUID.

dhcp-options-get-options dhcp-option

Lists the DHCP Options for the dhcp-option UUID.

DATABASE COMMANDS

These commands query and modify the contents of **ovsdb** tables. They are a slight abstraction of the **ovsdb** interface and as suchthey operate at a lower level than other **ovn-nbctl** commands.

Identifying Tables, Records, and Columns

Each of these commands has a *table* parameter to identify a table within the database. Many of them also take a *record* parameter that identifies a particular record within a table. The *record* parameter may be the UUID for a record, and many tables offer additional ways to identify records. Some commands also take *column* parameters that identify a particular field within the records in a table.

The following tables are currently defined:

Logical_Switch

An L2 logical switch. Records may be identified by name.

Logical_Switch_Port

A port within an L2 logical switch. Records may be identified by name.

ACL An ACL rule for a logical switch that points to it through its *acls* column.

Logical_Router

An L3 logical router. Records may be identified by name.

Logical_Router_Port

A port within an L3 logical router. Records may be identified by name.

Logical_Router_Static_Route

A static route belonging to an L3 logical router.

Address Set

An address set that can be used in ACLs.

Load_Balancer

A load balancer for a logical switch that points to it through its *load_balancer* column.

NAT A NAT rule for a Gateway router.

DHCP_Options

DHCP options.

NB Global

North bound global configurations.

Database Values

Each column in the database accepts a fixed type of data. The currently defined basic types, and their representations, are:

integer A decimal integer in the range -2**63 to 2**63-1, inclusive.

real A floating-point number.

Boolean

True or false, written **true** or **false**, respectively.

An arbitrary Unicode string, except that null bytes are not allowed. Quotes are optional for most strings that begin with an English letter or underscore and consist only of letters, underscores, hyphens, and periods. However, **true** and **false** and strings that match the syntax of UUIDs (see below) must be enclosed in double quotes to distinguish them from other basic types. When double quotes are used, the syntax is that of strings in JSON, e.g. backslashes may be used to escape special characters. The empty string must be represented as a pair of double quotes ("").

UUID Either a universally unique identifier in the style of RFC 4122, e.g. **f81d4fae-7dec-11d0-a765-00a0c91e6bf6**, or an @name defined by a **get** or **create** command within the same **ovn-nbctl** invocation.

Multiple values in a single column may be separated by spaces or a single comma. When multiple values are present, duplicates are not allowed, and order is not important. Conversely, some database columns can have an empty set of values, represented as [], and square brackets may optionally enclose other non-empty sets or single values as well.

A few database columns are "maps" of key-value pairs, where the key and the value are each some fixed database type. These are specified in the form *key=value*, where *key* and *value* follow the syntax for the column's key type and value type, respectively. When multiple pairs are present (separated by spaces or a comma), duplicate keys are not allowed, and again the order is not important. Duplicate values are allowed. An empty map is represented as {}. Curly braces may optionally enclose non-empty maps as well (but use quotes to prevent the shell from expanding **other-config={0=x,1=y}** into **other-config=0=x other-config=1=y**, which may not have the desired effect).

Database Command Syntax

[--if-exists] [--columns=column[,column]...] list table [record]...

Lists the data in each specified *record*. If no records are specified, lists all the records in *table*.

If **--columns** is specified, only the requested columns are listed, in the specified order. Otherwise, all columns are listed, in alphabetical order by column name.

Without — **if**—**exists**, it is an error if any specified *record* does not exist. With — **if**—**exists**, the command ignores any *record* that does not exist, without producing any output.

[--columns=column[,column]...] find table [column[:key]=value]...

Lists the data in each record in *table* whose *column* equals *value* or, if *key* is specified, whose *column* contains a *key* with the specified *value*. The following operators may be used where = is written in the syntax summary:

=!=<><=>=

Selects records in which *column*[:key] equals, does not equal, is less than, is greater than, is less than or equal to, or is greater than or equal to *value*, respectively.

Consider *column*[:key] and *value* as sets of elements. Identical sets are considered equal. Otherwise, if the sets have different numbers of elements, then the set with more elements is considered to be larger. Otherwise, consider a element from each set pairwise, in increasing order within each set. The first pair that differs determines the result. (For a column that contains key-value pairs, first all the keys are compared, and values are considered only if the two sets contain identical keys.)

{=} {!=}

Test for set equality or inequality, respectively.

{<=} Selects records in which column[:key] is a subset of value. For example, flood-vlans{<=}1,2 selects records in which the flood-vlans column is the empty set or contains 1 or 2 or both.</p>

Selects records in which column[:key] is a proper subset of value. For example, flood-vlans{<}1,2 selects records in which the flood-vlans column is the empty set or contains 1 or 2 but not both.</p>

{>=} {>}

Same as {<=} and {<}, respectively, except that the relationship is reversed. For example, **flood–vlans**{>=}1,2 selects records in which the **flood–vlans** column contains both 1 and 2.

For arithmetic operators (= != < > <= >=), when *key* is specified but a particular record's *column* does not contain *key*, the record is always omitted from the results. Thus, the condition **other-config:mtu!=1500** matches records that have a **mtu** key whose value is not 1500, but not those that lack an **mtu** key.

For the set operators, when *key* is specified but a particular record's *column* does not contain *key*, the comparison is done against an empty set. Thus, the condition **other–config:mtu{!=}1500** matches records that have a **mtu** key whose value is not 1500 and those that lack an **mtu** key.

Don't forget to escape < or > from interpretation by the shell.

If **—columns** is specified, only the requested columns are listed, in the specified order. Otherwise all columns are listed, in alphabetical order by column name.

The UUIDs shown for rows created in the same **ovn-nbctl** invocation will be wrong.

[--if-exists] [--id=@name] get table record [column[:key]]...

Prints the value of each specified *column* in the given *record* in *table*. For map columns, a *key* may optionally be specified, in which case the value associated with *key* in the column is printed, instead of the entire map.

Without ——**if**—**exists**, it is an error if *record* does not exist or *key* is specified, if *key* does not exist in *record*. With ——**if**—**exists**, a missing *record* yields no output and a missing *key* prints a blank line.

If @name is specified, then the UUID for record may be referred to by that name later in the same **ovn-nbctl** invocation in contexts where a UUID is expected.

Both ——id and the *column* arguments are optional, but usually at least one or the other should be specified. If both are omitted, then **get** has no effect except to verify that *record* exists in *table*.

--id and --if-exists cannot be used together.

[--if-exists] set table record column[:key]=value...

Sets the value of each specified *column* in the given *record* in *table* to *value*. For map columns, a *key* may optionally be specified, in which case the value associated with *key* in that column is changed (or added, if none exists), instead of the entire map.

Without --if-exists, it is an error if *record* does not exist. With --if-exists, this command does nothing if *record* does not exist.

[--if-exists] add table record column [key=]value...

Adds the specified value or key-value pair to *column* in *record* in *table*. If *column* is a map, then *key* is required, otherwise it is prohibited. If *key* already exists in a map column, then the current *value* is not replaced (use the **set** command to replace an existing value).

Without ——if—exists, it is an error if *record* does not exist. With ——if—exists, this command does nothing if *record* does not exist.

[--if-exists] remove table record column value...

[--if-exists] remove table record column key...

[—**if**—**exists**] **remov** *table record column key=value*... Removes the specified values or key-value pairs from *column* in *record* in *table*. The first form applies to columns that are not maps: each specified *value* is removed from the column. The second and third forms apply to map columns: if only a *key* is specified, then any key-value pair with the given *key* is removed, regardless of its value; if a *value* is given then a pair is removed only if both key and value match.

It is not an error if the column does not contain the specified key or value or pair.

Without ——if—exists, it is an error if *record* does not exist. With ——if—exists, this command does nothing if *record* does not exist.

[--if-exists] clear table record column...

Sets each *column* in *record* in *table* to the empty set or empty map, as appropriate. This command applies only to columns that are allowed to be empty.

Without ——**if**—**exists**, it is an error if *record* does not exist. With ——**if**—**exists**, this command does nothing if *record* does not exist.

[--id=@name] create table column[:key]=value...

Creates a new record in *table* and sets the initial values of each *column*. Columns not explicitly set will receive their default values. Outputs the UUID of the new row.

If @name is specified, then the UUID for the new row may be referred to by that name elsewhere in the same *(PN invocation in contexts where a UUID is expected. Such references may precede or follow the **create** command.

Caution (ovs-vsctl as exmaple)

Records in the Open vSwitch database are significant only when they can be reached directly or indirectly from the **Open_vSwitch** table. Except for records in the **QoS** or **Queue** tables, records that are not reachable from the **Open_vSwitch** table are automatically deleted from the database. This deletion happens immediately, without waiting for additional **ovs-vsctl** commands or other database activity. Thus, a **create** command must generally be accompanied by additional commands *within the same* **ovs-vsctl** *invocation* to add a chain of references to the newly created record from the top-level **Open_vSwitch** record. The **EXAMPLES** section gives some examples that show how to do this.

[--if-exists] destroy table record...

Deletes each specified *record* from *table*. Unless —**if**—**exists** is specified, each *records* must exist.

--all destroy table

Deletes all records from the table.

Caution (ovs-vsctl as exmaple)

The **destroy** command is only useful for records in the **QoS** or **Queue** tables. Records in other tables are automatically deleted from the database when they become unreachable from the **Open_vSwitch** table. This means that deleting the last reference to a record is sufficient for deleting the record itself. For records in these tables, **destroy** is silently ignored. See the **EXAMPLES** section below for more information.

wait-until table record [column[:key]=value]...

Waits until *table* contains a record named *record* whose *column* equals *value* or, if *key* is specified, whose *column* contains a *key* with the specified *value*. Any of the operators !=, <, >, <=, or >= may be substituted for = to test for inequality, less than, greater than, less than or equal to, or greater than or equal to, respectively. (Don't forget to escape < or > from interpretation by the shell.)

If no *column*[:*key*]=*value* arguments are given, this command waits only until *record* exists. If more than one such argument is given, the command waits until all of them are

satisfied.

Caution (ovs-vsctl as exmaple)

Usually **wait—until** should be placed at the beginning of a set of **ovs—vsctl** commands. For example, **wait—until bridge br0 — get bridge br0 datapath_id** waits until a bridge named **br0** is created, then prints its **datapath_id** column, whereas **get bridge br0 datapath_id** — **wait—until bridge br0** will abort if no bridge named **br0** exists when **ovs—vsctl** initially connects to the database.

Consider specifying —**timeout=0** along with —**wait—until**, to prevent **ovn—nbctl** from terminating after waiting only at most 5 seconds.

comment [arg]...

This command has no effect on behavior, but any database log record created by the command will include the command and its arguments.

SYNCHRONIZATION COMMANDS

Ordinarily, —wait=sb or —wait=hv only waits for changes by the current ovn—nbctl invocation to take effect. This means that, if none of the commands supplied to ovn—nbctl change the database, then the command does not wait at all. With the sync command, however, ovn—nbctl waits even for earlier changes to the database to propagate down to the southbound database or all of the OVN chassis, according to the argument to —wait.

REMOTE CONNECTIVITY COMMANDS

get-connection

Prints the configured connection(s).

del-connection

Deletes the configured connection(s).

set-connection *target...*

Sets the configured manager target or targets.

SSL CONFIGURATION COMMANDS

get-ssl Prints the SSL configuration.

del-ssl Deletes the current SSL configuration.

[—bootstrap] set—ssl private-key certificate ca-cert Sets the SSL configuration.

OPTIONS

```
--no-wait | --wait=none
--wait=sb
--wait=hv
```

These options control whether and how **ovn-nbctl** waits for the OVN system to become up-to-date with changes made in an **ovn-nbctl** invocation.

By default, or if **--no-wait** or **--wait=none**, **ovn-nbctl** exits immediately after confirming that changes have been committed to the northbound database, without waiting.

With —wait=sb, before ovn—nbctl exits, it waits for ovn—northd to bring the southbound database up-to-date with the northbound database updates.

With —wait=hv, before ovn—nbctl exits, it additionally waits for all OVN chassis (hypervisors and gateways) to become up-to-date with the northbound database updates. (This can become an indefinite wait if any chassis is malfunctioning.)

Ordinarily, —wait=sb or —wait=hv only waits for changes by the current ovn—nbctl invocation to take effect. This means that, if none of the commands supplied to ovn—nbctl change the database, then the command does not wait at all. Use the sync command to override this behavior.

--db database

The OVSDB database remote to contact. If the OVN_NB_DB environment variable is set, its value is used as the default. Otherwise, the default is unix:/usr/local/var/run/open-vswitch/ovnnb_db.sock, but this default is unlikely to be useful outside of single-machine OVN test environments.

LOGGING OPTIONS

$-\mathbf{v}[spec]$

--verbose=[spec]

Sets logging levels. Without any *spec*, sets the log level for every module and destination to **dbg**. Otherwise, *spec* is a list of words separated by spaces or commas or colons, up to one from each category below:

- A valid module name, as displayed by the **vlog/list** command on **ovs-appctl**(8), limits the log level change to the specified module.
- syslog, console, or file, to limit the log level change to only to the system log, to the console, or to a file, respectively. (If --detach is specified, the daemon closes its standard file descriptors, so logging to the console will have no effect.)

On Windows platform, **syslog** is accepted as a word and is only useful along with the **—syslog—target** option (the word has no effect otherwise).

• **off**, **emer**, **err**, **warn**, **info**, or **dbg**, to control the log level. Messages of the given severity or higher will be logged, and messages of lower severity will be filtered out. **off** filters out all messages. See **ovs-appctl**(8) for a definition of each log level.

Case is not significant within spec.

Regardless of the log levels set for **file**, logging to a file will not take place unless —**log-file** is also specified (see below).

For compatibility with older versions of OVS, any is accepted as a word but has no effect.

$-\mathbf{v}$

--verbose

Sets the maximum logging verbosity level, equivalent to **--verbose=dbg**.

- vPATTERN: destination: pattern

--verbose=PATTERN:destination:pattern

Sets the log pattern for *destination* to *pattern*. Refer to **ovs-appctl**(8) for a description of the valid syntax for *pattern*.

-vFACILITY: facility

--verbose=FACILITY:facility

Sets the RFC5424 facility of the log message. *facility* can be one of **kern**, **user**, **mail**, **daemon**, **auth**, **syslog**, **lpr**, **news**, **uucp**, **clock**, **ftp**, **ntp**, **audit**, **alert**, **clock2**, **local0**, **local1**, **local2**, **local3**, **local4**, **local5**, **local6** or **local7**. If this option is not specified, **daemon** is used as the default for the local system syslog and **local0** is used while sending a message to the target provided via the **--syslog-target** option.

--log-file[=file]

Enables logging to a file. If *file* is specified, then it is used as the exact name for the log file. The default log file name used if *file* is omitted is /usr/local/var/log/openvswitch/program.log.

--syslog-target=host:port

Send syslog messages to UDP *port* on *host*, in addition to the system syslog. The *host* must be a numerical IP address, not a hostname.

--syslog-method=method

Specify *method* as how syslog messages should be sent to syslog daemon. The following forms are supported:

- libc, to use the libc syslog() function. This is the default behavior. Downside of using this
 options is that libc adds fixed prefix to every message before it is actually sent to the syslog daemon over /dev/log UNIX domain socket.
- unix:file, to use a UNIX domain socket directly. It is possible to specify arbitrary message format with this option. However, rsyslogd 8.9 and older versions use hard coded parser function anyway that limits UNIX domain socket use. If you want to use arbitrary message format with older rsyslogd versions, then use UDP socket to localhost IP address instead.
- udp:ip:port, to use a UDP socket. With this method it is possible to use arbitrary message format also with older rsyslogd. When sending syslog messages over UDP socket extra precaution needs to be taken into account, for example, syslog daemon needs to be configured to listen on the specified UDP port, accidental iptables rules could be interfering with local syslog traffic and there are some security considerations that apply to UDP sockets, but do not apply to UNIX domain sockets.

PKI Options

PKI configuration is required to use SSL for the connection to the database.

−p privkey.pem

--private-key=privkey.pem

Specifies a PEM file containing the private key used as identity for outgoing SSL connections

-c cert.pem

--certificate=cert.pem

Specifies a PEM file containing a certificate that certifies the private key specified on **-p** or **--private-key** to be trustworthy. The certificate must be signed by the certificate authority (CA) that the peer in SSL connections will use to verify it.

-C cacert.pem

--ca-cert=cacert.pem

Specifies a PEM file containing the CA certificate for verifying certificates presented to this program by SSL peers. (This may be the same certificate that SSL peers use to verify the certificate specified on **–c** or **––certificate**, or it may be a different one, depending on the PKI design in use.)

-C none

--ca-cert=none

Disables verification of certificates presented by SSL peers. This introduces a security risk, because it means that certificates cannot be verified to be those of known trusted hosts.

--bootstrap-ca-cert=cacert.pem

When *cacert.pem* exists, this option has the same effect as $-\mathbf{C}$ or $--\mathbf{ca}-\mathbf{cert}$. If it does not exist, then the executable will attempt to obtain the CA certificate from the SSL peer on its first SSL connection and save it to the named PEM file. If it is successful, it will immediately drop the connection and reconnect, and from then on all SSL connections must be authenticated by a certificate signed by the CA certificate thus obtained.

This option exposes the SSL connection to a man-in-the-middle attack obtaining the initial CA certificate, but it may be useful for bootstrapping.

This option is only useful if the SSL peer sends its CA certificate as part of the SSL certificate chain. The SSL protocol does not require the server to send the CA certificate.

This option is mutually exclusive with **–**C and **––ca–cert**.

Other Options

-h

- --help Prints a brief help message to the console.
- $-\mathbf{V}$
- --version

Prints version information to the console.