

## NAME

ovn-northd – Open Virtual Network central control daemon

## SYNOPSIS

**ovn-northd** [*options*]

## DESCRIPTION

**ovn-northd** is a centralized daemon responsible for translating the high-level OVN configuration into logical configuration consumable by daemons such as **ovn-controller**. It translates the logical network configuration in terms of conventional network concepts, taken from the OVN Northbound Database (see **ovn-nb**(5)), into logical datapath flows in the OVN Southbound Database (see **ovn-sb**(5)) below it.

## CONFIGURATION

**ovn-northd** requires a connection to the Northbound and Southbound databases. The default is **db.sock** in the local Open vSwitch's "run" directory. This may be overridden with the following commands:

- **--ovnnb-db=database**  
The database containing the OVN Northbound Database.
- **--ovsdb-db=database**  
The database containing the OVN Southbound Database.

The *database* argument must take one of the following forms:

- **ssl:ip:port**  
The specified SSL *port* on the host at the given *ip*, which must be expressed as an IP address (not a DNS name) in IPv4 or IPv6 address format. If *ip* is an IPv6 address, then wrap *ip* with square brackets, e.g.: **ssl:::1]:6640**. The **--private-key**, **--certificate**, and **--ca-cert** options are mandatory when this form is used.
- **tcp:ip:port**  
Connect to the given TCP *port* on *ip*, where *ip* can be IPv4 or IPv6 address. If *ip* is an IPv6 address, then wrap *ip* with square brackets, e.g.: **tcp:::1]:6640**.
- **unix:file**  
On POSIX, connect to the Unix domain server socket named *file*.  
On Windows, connect to a localhost TCP port whose value is written in *file*.

## RUNTIME MANAGEMENT COMMANDS

**ovs-appctl** can send commands to a running **ovn-northd** process. The currently supported commands are described below.

- exit** Causes **ovn-northd** to gracefully terminate.

## LOGICAL FLOW TABLE STRUCTURE

One of the main purposes of **ovn-northd** is to populate the **Logical\_Flow** table in the **OVN\_Southbound** database. This section describes how **ovn-northd** does this for logical datapaths.

### Ingress Table 0: Admission Control and Ingress Port Security

Ingress table 0 contains these logical flows:

- Priority 100 flows to drop packets with VLAN tags or multicast Ethernet source addresses.
- Priority 50 flows that implement ingress port security for each enabled logical port. For logical ports on which port security is enabled, these match the **inport** and the valid **eth.src** address(es) and advance only those packets to the next flow table. For logical ports on which port security is not enabled, these advance all packets that match the **inport**.

There are no flows for disabled logical ports because the default-drop behavior of logical flow tables causes packets that ingress from them to be dropped.

### Ingress table 1: from-lport ACLs

Logical flows in this table closely reproduce those in the **ACL** table in the **OVN\_Northbound** database for the **from-lport** direction. **allow** and **allow-related** ACLs translate into logical flows with the **next;** action, others to **drop;**. The **priority** values from the **ACL** table are used directly.

Ingress table 1 also contains a priority 0 flow with action **next;**, so that ACLs allow packets by default.

### Ingress Table 2: Destination Lookup

This table implements switching behavior. It contains these logical flows:

- A priority-100 flow that outputs all packets with an Ethernet broadcast or multicast **eth.dst** to the **MC\_FLOOD** multicast group, which **ovn-northd** populates with all enabled logical ports.
- One priority-50 flow that matches each known Ethernet address against **eth.dst** and outputs the packet to the single associated output port.
- One priority-0 fallback flow that matches all packets and outputs them to the **MC\_UNKNOWN** multicast group, which **ovn-northd** populates with all enabled logical ports that accept unknown destination packets. As a small optimization, if no logical ports accept unknown destination packets, **ovn-northd** omits this multicast group and logical flow.

### Egress Table 0: to-lport ACLs

This is similar to ingress table 1 except for **to-lport** ACLs.

### Egress Table 1: Egress Port Security

This is similar to the ingress port security logic in ingress table 0, but with important differences. Most obviously, **outport** and **eth.dst** are checked instead of **inport** and **eth.src**. Second, packets directed to broadcast or multicast **eth.dst** are always accepted instead of being subject to the port security rules; this is implemented through a priority-100 flow that matches on **eth.dst[40]** with action **output;**. Finally, to ensure that even broadcast and multicast packets are not delivered to disabled logical ports, a priority-150 flow for each disabled logical **outport** overrides the priority-100 flow with a **drop;** action.