

Container Networking Solutions

Agenda


• Introduction

- Container Network Functions & Interfaces
- Limitations
- Container Interface Classifier
- Community Solutions
- Our Proposed Solution
- Consolidation

• Functionality Offload

- P4 Sample -Kubeproxy
- P4 Connection tracking
- P4 L4-L7 Classifier


Container Networking



Current State of Deployment




Need for Scale
CSP Container
Scale Needed
~10K



Device performance
Throughput,
Programmability,
Resource Scaling
has increased
significantly



Effective utilization
of the system
resources (cores,
memory, network) to
improve
performance and
packing of containers



**Ideal
Container
Network
Solution**



5. Existing interfaces for containers

Shared (Pod Interfaces)

- Stacked netdevs on a PCIe PF netdev
- Assigned to container namespaces
- Examples : MACVLAN, IPVLAN, bridge

Dedicated (Pod Interfaces)

- SR-IOV VFs
- Too heavy
 - Separate PCIe config space
 - HW based packet replication for broadcast, multicast – higher PCI BW utilization

Why Hardware Accelerate ?

- End to end Maximize throughput
 - Avoid the SW long pole which limits how much a Server Pod can handle.
- End to end Native Scale out.
 - By Reduce Latency and Jitter
 - Dedicated resources takes away the need for OS to schedule on a shared resource.....OS overhead for managing resources is gone.
 - Cpu scheduling, memory management etc
- Security & Isolation
 - Queue level isolation.

Assignable Container interfaces using X-IOV

(S-IOV)

Hardware Assisted Virtualization

Highly scalable and high- performance sharing of I/O devices across isolated domains


Assignable device Interfaces =>User Container Interfaces

Platform Scalability using PASID

Support Virtual Device Composition


RID and PASID identifies the address space associated with the request

ADI Memory Mapped regions



4. Container Interface Classifier: Solution

- L4-L7 Classifier and forwarder in the HW
- This extends the HW Offloaded vSwitch Classification End Point to the Container Interface.
- Option1: HW Offload the classifier & forwarder
 - AF_XDP raw_socket bound to a HW vPort/QP through Side band filters.
 - Provide inline filters to be added in HW as part of TX packet.
 - ATR style in ADQ
- Option2: HW Offload the Classifier
 - Provide a meta data classification hint to kernel/user with a packet.
 - Flow mark or a 32bit hash value based on L4-L7 fields.



3. Node: Load Balancer, DNAT, CT

Existing Solutions

- **Kube-proxy**
 - Kernel Netfilters - Not performant
 - Iptables $O(n)$ chains proportional to size of cluster, in-place rule modifications not possible.
 - IPVS $O(1)$ - hash ipset but do not work well with other services requiring iptables for filtering
 - Kernel with eBF/XDP - Accelerated
- **Connection Tracking**
 - Robust to syn floods but limited by max size
 - No flexibility, fixed hash algorithm and field selection for hash
- **Kernel** Overall not very flexible, latencies due to irq processing, context switching, slow API configuration interfaces

Community's Approach - eBPF/XDP

Benefits - Performant than the kernel

Designed as an alternative to DPDK.

Flexibility, code injected into the kernel


Ability to reload programs on-the-fly

Network Functions - Network Policy, Encryption, Load Balancer, Firewall, Monitoring etc

Plugins - Cilium, Calico, Katran (facebook), etc

Limitations -

- Not HW offloadable.
- Limited by performance and scalability.
- Kernel/user space context switching
- General purpose CPUs and Memory architecture is not ideal for Deep table Lookup
- Scale in CPU usage/CPU cache with load



eBPF Implementation – Cilium

Cilium does connect-time load balancing by hooking into the BPF & XDP/TC hook on the receive. When a program tries to connect to a Kubernetes service, Cilium intercepts the connection attempt, load balances with DNAT's to directly connect to the backend pod's IP instead


Throughput in queries/s – test run with fortio and nginx. 2 clients, 2 services

No XDP, 2 CLV interfaces			
Fortio + Nginx	Client 1	Client 2	Total
100B	47578.6	50411.5	97990.1
1500B	46174.4	49990.4	96164.8
XDP, 1 CLV interface			
Fortio + Nginx	Client 1	Client 2	Total
100B	64038.7	53970	118008.7
1500B	49623.5	48369.7	97993.2

No-XDP and XDP performed similar with 1000 connections. CPU consumption in both reaches to >8 cores with more sessions

XDP path today does not take advantage of hardware XDP_REDIRECT queue designed to send packet to another interface, so no significant performance gains

Our Approach – The H/W Datapath



Programmable MAT tables. Contract between control plane and data plane for runtime control


Device capability defined by architecture. Eg wildcard match -TCAM, Exact match-DRAM. Compiler responsible for mapping.

Parallel lookups, conditional actions & atomicity


Features are defined in the software. Faster introduction, verification, test and deployment.

Programmer defined 1) parser (parse graph); headers (inner/outer L2,L3,L4, app) and orders 2) how the output packet will look on the wire

Counters, meters, stateful registers, hash functions, ALU, TTLs and counters per entry, PRE



Node LB Data Plane P4



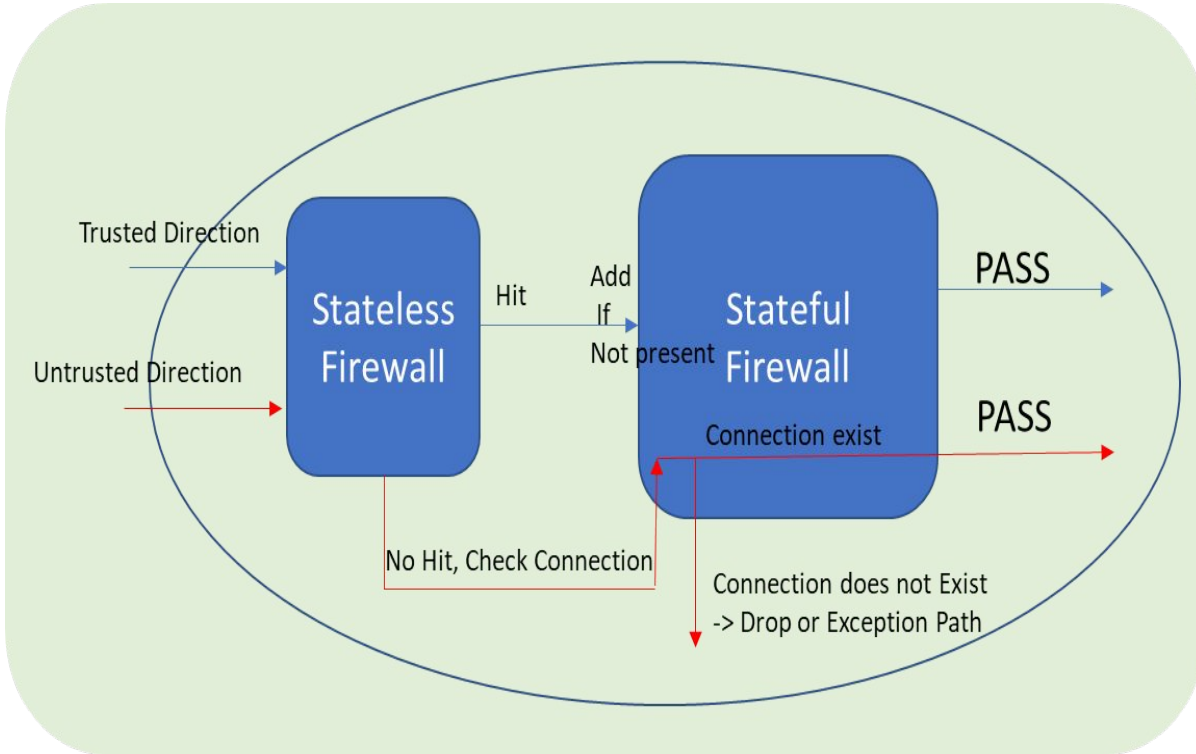
7/16/2019

```

1 // Proxy LB on cluster service IP to an endpoint on a POD
2 extern ActionSelector{
3     /// Construct a selection table for a given ActionProfile.
4     ActionSelector(ActionProfile action_profile,
5                     Hash<_> hash,
6                     SelectorMode_t mode,
7                     bit<32> max_group_size,
8                     bit<32> num_groups);
9
10    ActionSelector(bit<32> size, Hash<_> hash, SelectorMode_t mode);
11 }
12
13
14 control simple_lb(inout headers hdr,
15                  inout metadata meta,
16                  switch_uint32_t lb_table_size,
17                  inout standard_metadata_t standard_metadata) {
18
19     //Chose an extern hash or PNA hash
20     Hash<switch_uint32_t>(HashAlgorithm_t.CRC32) selector_hash;
21
22     ActionSelector(
23         1024, selector_hash, SelectorMode_t.FAIR) pod_selector;
24
25     // Pick an entry and apply DNAT
26     action set_nhop(bit<48> pod_dmac, bit<32> pod_ipv4, bit<9> port) {
27         hdr.ethernet.dstAddr = pod_dmac;
28         hdr.ipv4.dstAddr = pod_ipv4;
29         hdr.tcp.dstPort = port;
30         standard_metadata.egress_spec = port;
31     }
32
33     // hash to use 5-tuple
34     // can be tcp, udp, sctp
35     table lb {
36         key = {
37             hdr.ipv4.dstAddr : exact;
38             hdr.tcp.dstPort : exact;
39             hdr.ipv4.srcAddr : selector;
40             hdr.ipv4.dstAddr : selector;
41             hdr.ipv4.tcp.srcPort : selector;
42             hdr.ipv4.tcp.dstPort : selector;
43             hdr.ipv4.protocol : selector
44         }
45
46         actions = {
47             NoAction;
48         }
49     }
50 }


```


CT P4 Data Plane



```
C: > Users > asinghai > OneDrive - Intel Corporation > Documents > MEV > CT_mev.p4
256
257 struct ct_aging_table_hit_params_t {
258     FlowId_t flow_id;
259     ExpireTimeSelection_t expire_time;
260 }
261
262 action ct_aging_table_hit ( // See "Note 6"
263     @per_entry_state in FlowId_t flow_id, // See "Note 0"
264     @per_entry_state inout ExpireTimeSelection_t expire_time)
265 {
266     // my_flow_id = flow_id; // See "Note 1"
267     if (modify_expire_soon_on_hit) {
268         expire_time = new_expire_time_selection; // called TIME_SEL in MEV
269         modify_entry_to_expire_soon(); // See "Note 2"
270     } else if (update_expire_time) { // See "Note 3"
271         expire_time = new_expire_time_selection; // called TIME_SEL in MEV
272         restart_expire_timer(); // sweep_count = 0
273     } else {
274         restart_expire_timer(); // sweep_count = 0
275     }
276 }
277
278 action ct_aging_table_miss() {
279     FlowId_t my_flow_id; // See "Note 7"
280     bool add_succeeded;
281     if (add_on_miss) { // See "Note 4"
282         my_flow_id = allocate_flow_id(); // See "Note 1"
283         add_succeeded =
284             add_entry("ct_aging_table_hit", // name of action
285                     (ct_aging_table_hit_params_t)
286                     {flow_id = my_flow_id,
287                     expire_time = new_expire_time_selection});
288         // add_entry() initializes the new entry as if
289         // restart_expire_timer() had been called on it,
290         // e.g. for MEV, initialize a hidden 'sweep count' to 0
291     }
292 }
293
294 // LEM table
295 table ct_aging_table {
296     key = {
297         // P4 developer gets to select the key fields they want,
298         // e.g. This is a 5 tuple plus a field like zone id, used to guarantee
299         // that entries in different IP private address domains are
300         // unique.
301     }
302     actions = {
303         // @tableonly and @defaultonly are standard annotations
304         // defined in the P4 16 language specification.
```

L4-L7 P4 Classifier



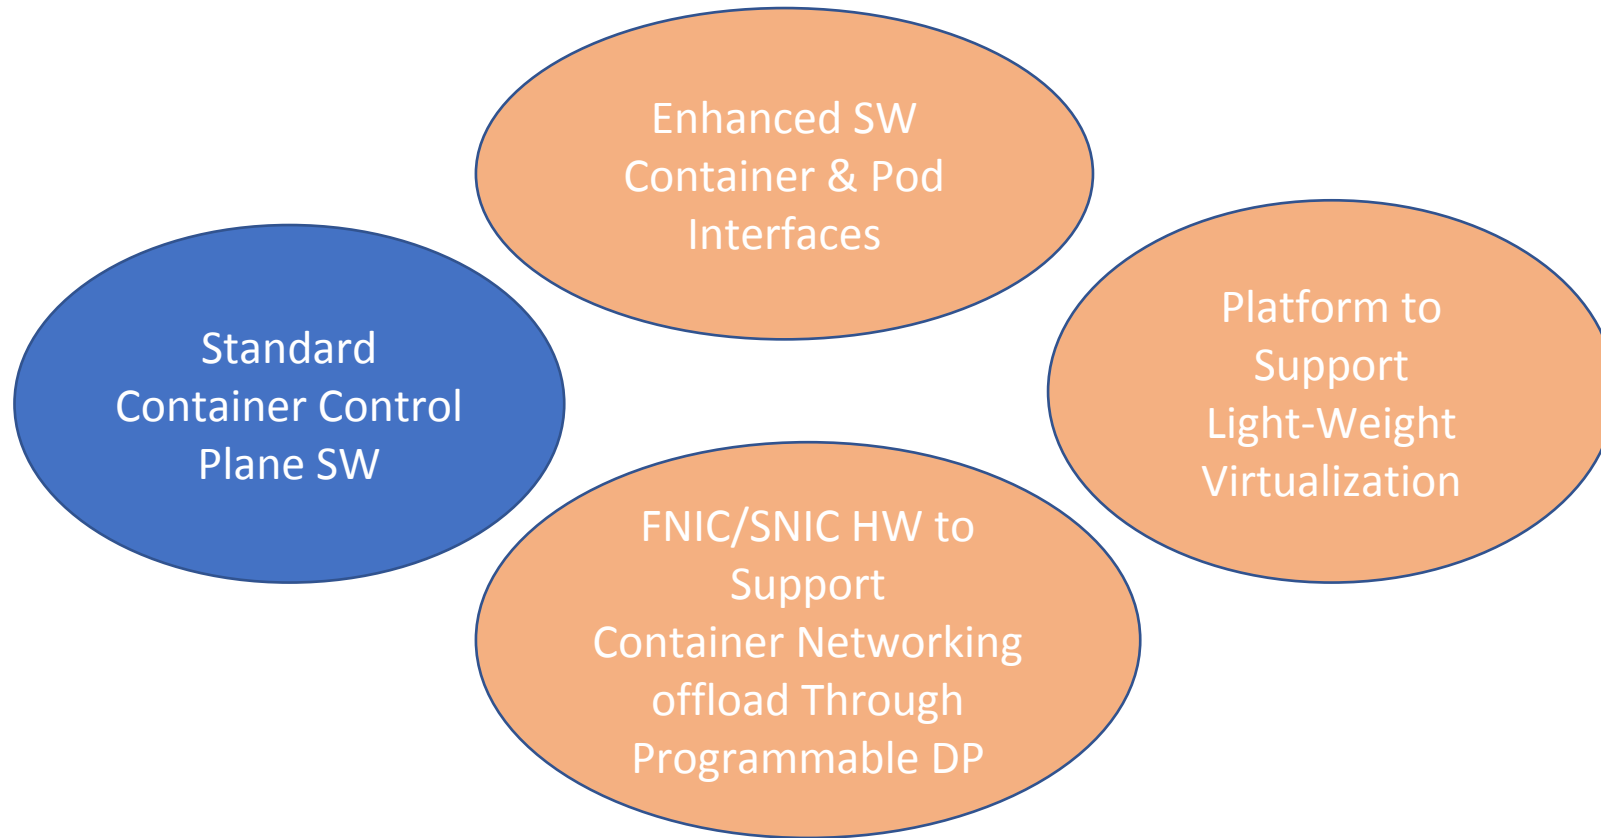
```
1 control L4_L7_classifier(inout headers hdr,  
2     inout metadata meta,  
3     switch_uint32_t lb_table_size,  
4     inout standard_metadata_t standard_metadata) {  
5  
6     // Forward to a Container Interface  
7     action set_queue(bit<16> queue_id) {  
8         standard_metadata.egress_spec = queue_id;  
9     }  
10  
11     // Set a flow ID meta data  
12     action set_flow_id(bit<32> mark) {  
13         standard_metadata.flow_id = mark;  
14     }  
15  
16     // exact match using n-tuple  
17     // can be tcp, udp, sctp  
18     table lb {  
19         key = {  
20             hdr.ipv4.dstAddr : exact;  
21             hdr.udp.dstPort : exact;  
22             hdr.udp.srcPort : exact;  
23             hdr.ipv4.srcAddr : exact;  
24             hdr.udp.quic.cid : exact;  
25         }  
26  
27         actions = {  
28             NoAction;  
29             set_queue;  
30             flow_count.count();  
31             set_flow_id;  
32         }  
33  
34         const default_action = Drop;  
35         counters = flow_counter;  
36         size = lb_table_size; // can be configured  
37         implementation = pod_selector;  
38     }  
39 }  
40  
41
```

Opens

- Not every eBPF program can be HW offloaded as is. We are looking at all use cases.
- We would like to get community support in converting some well defined XDP implementations to p4 programs.
- P4 extensions or externs is an option for complete packet transformations like Crypto, checksum, packet replication etc.

Conclusion:

Components for Native Scale out of Container Networking



Contacts

- P4 Code will be on github soon...
- Please email for more info...
- Contacts: Anjali Singhai anjali.singhai@intel.com
Nupur Jain nupur.jain@intel.com
- Intel Container Networking Team:

Anjali Singhai, Nupur Jain, Amritha Nambiar, Pawel Szymanski, Shaopeng He, Phani Burra,	Dan Daly, Yadong Li, Sridhar Samudralla, Kiran Patil, Liang Cunming , Edwin Verplanke

eBPF Implementation Characterization – Cilium

Stack trace (NO XDP)

vmlinux	615.317s
nginx	55.704s
[Unknown]	24.101s
[Outside any known module]	24.101s
cls_bpf_classify	21.866s
↖ __tcf_classify ← tcf_classify_in	1.173s
↖ __tcf_classify ← tcf_classify ←	1.042s
[Unknown]	0.020s
↖ func@0x7a5bd4 ← func@0x79	0s
func@0x7a0570	0s
func@0x3527ee1	0s
ice.ko	23.259s
ip_tables.ko	20.728s
libc-2.13.so	11.371s
amplxe-perf	8.445s
cls_bpf.ko	8.339s
containerd-shim	8.209s

Stack trace (XDP)

[Unknown]	28.626s	52,704,500,000
[Outside any known module]	28.626s	52,704,500,000
cls_bpf_classify	17.400s	29,923,000,000
↖ bpf_prog_run_xdp ←	8.194s	21,148,500,000
↖ __tcf_classify ← tcf_cl	1.463s	609,500,000
↖ __tcf_classify ← tcf_cl	1.273s	621,000,000
↖ ice_napi_poll ← napi_	0.236s	322,000,000
[Unknown]	0.035s	23,000,000
func@0x79fe60	0.015s	11,500,000
↖ func@0x7a19d5 ← fur	0.005s	0
func@0x3527ee1	0.005s	34,500,000
↖ func@0x7a0570 ← fur	0s	11,500,000
ice.ko	25.519s	26,553,500,000
ip_tables.ko	17.896s	32,798,000,000
libc-2.13.so	11.096s	22,482,500,000
containerd-shim	7.943s	12,788,000,000
cls_bpf.ko	7.367s	8,533,000,000

- XDP vs No XDP, CPU utilization is quite similar
- Benefits of XDP is being able to bypass the kernel stack in case of redirect.
- Redirect to external port requires dedicate HW Redirect TX queue.
- XDP Benefits can be derived from dedicated HW resources.

eBPF Implementation Characterization – Cilium - cont



Cilium chains in iptables

```
# Generated by iptables-save v1.6.1 on Tue Dec 1 13:22:26 2020
*raw
:PREROUTING ACCEPT [339745057:40541787768]
:OUTPUT ACCEPT [1966765:7841806195]
:CILIUM_OUTPUT_raw - [0:0]
:CILIUM_PRE_raw - [0:0]
-A PREROUTING -m comment --comment "cilium-feeder: CILIUM_PRE_raw" -j CILIUM_PRE_raw
-A OUTPUT -m comment --comment "cilium-feeder: CILIUM_OUTPUT_raw" -j CILIUM_OUTPUT_raw
-A CILIUM_OUTPUT_raw -o lxc+ -m mark --mark 0xa00/0xffffffff -m comment --comment "cilium: NOTRACK for proxy return traff
-A CILIUM_OUTPUT_raw -o cilium_host -m mark --mark 0xa00/0xffffffff -m comment --comment "cilium: NOTRACK for proxy retur
-A CILIUM_PRE_raw -m mark --mark 0x200/0xf00 -m comment --comment "cilium: NOTRACK for proxy traffic" -j NOTRACK
COMMIT
# Completed on Tue Dec 1 13:22:26 2020
# Generated by iptables-save v1.6.1 on Tue Dec 1 13:22:26 2020
*mangle
:PREROUTING ACCEPT [339745170:40541800199]
:INPUT ACCEPT [1689410:7607891639]
:FORWARD ACCEPT [338055760:32933908560]
:OUTPUT ACCEPT [1966765:7841806195]
:POSTROUTING ACCEPT [340022530:40775715026]
:CILIUM_POST_mangle - [0:0]
:CILIUM_PRE_mangle - [0:0]
-A PREROUTING -m comment --comment "cilium-feeder: CILIUM_PRE_mangle" -j CILIUM_PRE_mangle
-A POSTROUTING -m comment --comment "cilium-feeder: CILIUM_POST_mangle" -j CILIUM_POST_mangle
-A CILIUM_PRE_mangle -m socket --transparent -m comment --comment "cilium: any->pod redirect proxied traffic to host prox
-A CILIUM_PRE_mangle -p tcp -m mark --mark 0xd7ae0200 -m comment --comment "cilium: TPROXY to host cilium-dns-egress prox
-A CILIUM_PRE_mangle -p udp -m mark --mark 0xd7ae0200 -m comment --comment "cilium: TPROXY to host cilium-dns-egress prox
COMMIT
# Completed on Tue Dec 1 13:22:26 2020
# Generated by iptables-save v1.6.1 on Tue Dec 1 13:22:26 2020
*filter
:INPUT ACCEPT [1689410:7607891639]
:FORWARD DROP [0:0]
:OUTPUT ACCEPT [1966765:7841806195]
:CILIUM_FORWARD - [0:0]
:CILIUM_INPUT - [0:0]
:CILIUM_OUTPUT - [0:0]
-A INPUT -m comment --comment "cilium-feeder: CILIUM_INPUT" -j CILIUM_INPUT
-A FORWARD -m comment --comment "cilium-feeder: CILIUM_FORWARD" -j CILIUM_FORWARD
-A OUTPUT -m comment --comment "cilium-feeder: CILIUM_OUTPUT" -j CILIUM_OUTPUT
-A CILIUM_FORWARD -o cilium_host -m comment --comment "cilium: any->cluster on cilium_host forward accept" -j ACCEPT
-A CILIUM_FORWARD -i cilium_host -m comment --comment "cilium: cluster->any on cilium_host forward accept (nodeport)" -j
-A CILIUM_FORWARD -i lxc+ -m comment --comment "cilium: cluster->any on lxc+ forward accept" -j ACCEPT
-A CILIUM_FORWARD -i cilium_net -m comment --comment "cilium: cluster->any on cilium_net forward accept (nodeport)" -j AC
-A CILIUM_INPUT -m mark --mark 0x200/0xf00 -m comment --comment "cilium: ACCEPT for proxy traffic" -j ACCEPT
-A CILIUM_OUTPUT -m mark --mark 0xa00/0xffffffff -m comment --comment "cilium: ACCEPT for proxy return traffic" -j ACCEPT
-A CILIUM_OUTPUT -m mark ! --mark 0xe00/0xf00 -m mark ! --mark 0xd00/0xf00 -m mark ! --mark 0xa00/0xe00 -m comment --comm
COMMIT
# Completed on Tue Dec 1 13:22:26 2020
# Generated by iptables-save v1.6.1 on Tue Dec 1 13:22:26 2020
*nat
:PREROUTING ACCEPT [0:0]
:INPUT ACCEPT [0:0]
```

Cilium's own conntrack table

```
TCP IN 12.91.212.200:37386 -> 10.0.0.241:80 expires=17272887 RxBPackets=18796 RxBBytes=2057656 RxFlagsSeen=0x1b LastRxReport=17251794 TxPackets=18915 TxBytes=4544990 TxFlagsSeen=
TCP OUT 13.91.212.200:33224 -> 10.0.0.55:80 expires=17275420 RxBPackets=83205 RxBBytes=19555155 RxFlagsSeen=0x1a LastRxReport=17254320 TxPackets=80142 TxBytes=9897097 TxFlagsSeen=
TCP OUT 12.91.212.202:30007 -> 12.91.212.200:36342 service expires=17272760 RxBPackets=0 RxBBytes=6 RxFlagsSeen=0x00 LastRxReport=0 TxPackets=0 TxBytes=0 TxFlagsSeen=0x1f LastTxR
TCP OUT 12.91.212.202:38900 -> 12.91.212.200:30008 expires=17274846 RxBPackets=2 RxBBytes=120 RxFlagsSeen=0x14 LastRxReport=17253753 TxPackets=1 TxBytes=74 TxFlagsSeen=0x02 LastT
TCP OUT 12.91.212.200:43362 -> 10.0.0.246:80 expires=17275288 RxBPackets=141178 RxBBytes=33482100 RxFlagsSeen=0x1a LastRxReport=17254194 TxPackets=137700 TxBytes=17005280 TxFlagsS
TCP IN 13.91.212.200:49604 -> 10.0.0.55:80 expires=17271556 RxBPackets=90047 RxBBytes=9859771 RxFlagsSeen=0x1b LastRxReport=17250463 TxPackets=91583 TxBytes=21846155 TxFlagsSeen=
TCP IN 13.91.212.200:55058 -> 10.0.0.55:80 expires=17273171 RxBPackets=52349 RxBBytes=5757922 RxFlagsSeen=0x1b LastRxReport=17252078 TxPackets=53989 TxBytes=49437539 TxFlagsSeen=
TCP OUT 12.91.212.200:43004 -> 10.0.0.241:80 expires=17275006 RxBPackets=1016 RxBBytes=227571 RxFlagsSeen=0x1a LastRxReport=17253912 TxPackets=923 TxBytes=113493 TxFlagsSeen=0x1e
TCP IN 12.91.212.200:44964 -> 10.0.0.246:80 expires=17275288 RxBPackets=140974 RxBBytes=17409857 RxFlagsSeen=0x1e LastRxReport=17254195 TxPackets=145276 TxBytes=34327857 TxFlagsS
TCP OUT 13.91.212.202:48040 -> 13.91.212.200:30008 expires=17274814 RxBPackets=2 RxBBytes=120 RxFlagsSeen=0x14 LastRxReport=17253721 TxPackets=1 TxBytes=74 TxFlagsSeen=0x02 LastT
TCP IN 12.91.212.200:44096 -> 10.0.0.241:80 expires=17275421 RxBPackets=102519 RxBBytes=12660607 RxFlagsSeen=0x1e LastRxReport=17254328 TxPackets=104458 TxBytes=24884841 TxFlagsS
TCP IN 12.91.212.200:43468 -> 10.0.0.246:80 expires=17275288 RxBPackets=130334 RxBBytes=16095587 RxFlagsSeen=0x1e LastRxReport=17254195 TxPackets=133864 TxBytes=31706643 TxFlagsS
TCP OUT 12.91.212.202:30007 -> 12.91.212.200:42992 service expires=17275006 RxBPackets=0 RxBBytes=6 RxFlagsSeen=0x00 LastRxReport=0 TxPackets=0 TxBytes=0 TxFlagsSeen=0x1e LastTxR
TCP OUT 13.91.212.202:30008 -> 13.91.212.200:54258 service expires=17272887 RxBPackets=0 RxBBytes=9 RxFlagsSeen=0x00 LastRxReport=0 TxPackets=0 TxBytes=0 TxFlagsSeen=0x1b LastTxR
TCP OUT 12.91.212.202:54410 -> 12.91.212.200:30007 expires=17274731 RxBPackets=2 RxBBytes=120 RxFlagsSeen=0x14 LastRxReport=17253638 TxPackets=1 TxBytes=74 TxFlagsSeen=0x02 LastT
TCP OUT 12.91.212.202:54426 -> 12.91.212.200:30007 expires=17274731 RxBPackets=2 RxBBytes=120 RxFlagsSeen=0x14 LastRxReport=17253638 TxPackets=1 TxBytes=74 TxFlagsSeen=0x02 LastT
TCP OUT 12.91.212.202:30007 -> 12.91.212.200:35150 service expires=17272401 RxBPackets=0 RxBBytes=8 RxFlagsSeen=0x00 LastRxReport=0 TxPackets=0 TxBytes=0 TxFlagsSeen=0x1f LastTxR
TCP IN 12.91.212.200:42930 -> 10.0.0.246:80 expires=17275006 RxBPackets=1174 RxBBytes=144442 RxFlagsSeen=0x1e LastRxReport=17253913 TxPackets=1292 TxBytes=289662 TxFlagsSeen=0x1a
TCP IN 13.91.212.200:33946 -> 10.0.0.55:80 expires=17275768 RxBPackets=24996 RxBBytes=3086447 RxFlagsSeen=0x1a LastRxReport=17254672 TxPackets=25670 TxBytes=6078920 TxFlagsSeen=0
TCP OUT 12.91.212.202:30007 -> 12.91.212.200:38580 service expires=17273496 RxBPackets=0 RxBBytes=6 RxFlagsSeen=0x00 LastRxReport=0 TxPackets=0 TxBytes=0 TxFlagsSeen=0x1b LastTxR
TCP IN 13.91.212.200:34256 -> 10.0.0.55:80 expires=17275769 RxBPackets=21454 RxBBytes=2649240 RxFlagsSeen=0x1a LastRxReport=17254675 TxPackets=22597 TxBytes=5255183 TxFlagsSeen=0
TCP OUT 13.91.212.200:33150 -> 10.0.0.55:80 expires=17275420 RxBPackets=80644 RxBBytes=19005294 RxFlagsSeen=0x1a LastRxReport=17254316 TxPackets=77972 TxBytes=9629102 TxFlagsSeen=
TCP OUT 12.91.212.202:30007 -> 12.91.212.200:44004 service expires=17275421 RxBPackets=0 RxBBytes=8 RxFlagsSeen=0x00 LastRxReport=0 TxPackets=0 TxBytes=0 TxFlagsSeen=0x1e LastTxR
TCP OUT 13.91.212.202:30008 -> 13.91.212.200:55654 service expires=17273480 RxBPackets=0 RxBBytes=9 RxFlagsSeen=0x00 LastRxReport=0 TxPackets=0 TxBytes=0 TxFlagsSeen=0x1b LastTxR
TCP OUT 12.91.212.202:55190 -> 12.91.212.200:30007 expires=17274765 RxBPackets=2 RxBBytes=120 RxFlagsSeen=0x14 LastRxReport=17253672 TxPackets=1 TxBytes=74 TxFlagsSeen=0x02 LastT
TCP OUT 12.91.212.202:39914 -> 12.91.212.200:30008 expires=17274929 RxBPackets=2 RxBBytes=120 RxFlagsSeen=0x14 LastRxReport=17253836 TxPackets=1 TxBytes=74 TxFlagsSeen=0x02 LastT
TCP OUT 12.91.212.202:38932 -> 12.91.212.200:30008 expires=17274846 RxBPackets=2 RxBBytes=120 RxFlagsSeen=0x14 LastRxReport=17253753 TxPackets=1 TxBytes=74 TxFlagsSeen=0x02 LastT
TCP OUT 12.91.212.202:38888 -> 12.91.212.200:30008 expires=17274846 RxBPackets=2 RxBBytes=120 RxFlagsSeen=0x14 LastRxReport=17253753 TxPackets=1 TxBytes=74 TxFlagsSeen=0x02 LastT
TCP OUT 12.91.212.202:30007 -> 12.91.212.200:44106 service expires=17275421 RxBPackets=0 RxBBytes=8 RxFlagsSeen=0x00 LastRxReport=0 TxPackets=0 TxBytes=0 TxFlagsSeen=0x1e LastTxR
TCP IN 12.91.212.200:44194 -> 10.0.0.246:80 expires=17275421 RxBPackets=115638 RxBBytes=14280861 RxFlagsSeen=0x1e LastRxReport=17254328 TxPackets=118303 TxBytes=28101171 TxFlagsS
TCP OUT 12.91.212.202:30007 -> 12.91.212.200:42722 service expires=17274983 RxBPackets=0 RxBBytes=8 RxFlagsSeen=0x00 LastRxReport=0 TxPackets=0 TxBytes=0 TxFlagsSeen=0x1e LastTxR
```

Q&A